

Deep Learning

Optimizers



Gananath Bhuyan
Assistant Professor

School of Computer Engineering
KIIT DU, Bhubaneswar

What are Optimizers in Deep Learning?

- In deep learning, an optimizer is a crucial element that fine-tunes a neural network's parameters during training. Its primary role is to minimize the model's error or loss function, enhancing performance.
- These specialized algorithms facilitate the learning process of neural networks by iteratively refining the weights and biases based on the feedback received from the data.
- Eg: **Stochastic Gradient Descent (SGD)**, **Mini-Batch Gradient Descent**, **Adam**, **AdaDelta**, and **RMSprop**, etc., each equipped with distinct update rules, learning rates, and momentum strategies.

Choosing the Right Optimizer

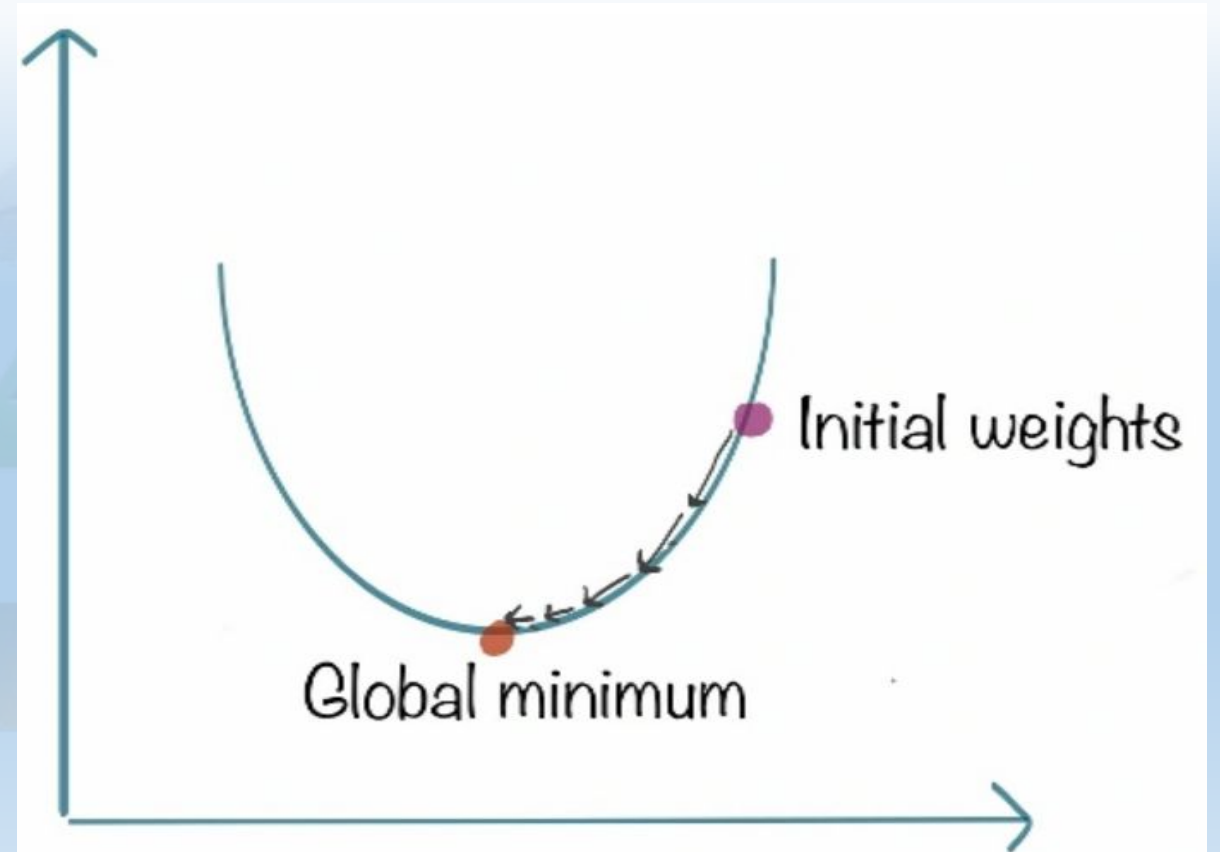
- **Epoch** – The number of times the algorithm runs on the whole training dataset.
- **Sample** – A single row of a dataset.
- **Batch** – It denotes the number of samples to be taken to for updating the model parameters.
- **Learning rate** – It is a parameter that provides the model a scale of how much model weights should be updated.
- **Cost Function/Loss Function** – A cost function is used to calculate the cost, which is the difference between the predicted value and the actual value.
- **Weights/ Bias** – The learnable parameters in a model that controls the signal between two neurons.

Gradient Descent

- This optimization algorithm uses calculus to consistently modify the values and achieve the local minimum.
1. Initialize Coefficients: Start with initial coefficients.
 2. Evaluate Cost: Calculate the cost associated with these coefficients.
 3. Search for Lower Cost: Look for a cost value lower than the current one.
 4. Update Coefficients: Move towards the lower cost by updating the coefficients' values.
 5. Repeat Process: Continue this process iteratively.
 6. Reach Local Minimum: Stop when a local minimum is reached, where further cost reduction is not possible.

Gradient Descent

- It is expensive to calculate the gradients if the size of the data is huge.
- Gradient descent works well for convex functions, but it doesn't know how far to travel along the gradient for nonconvex functions.



Stochastic Gradient Descent

Back

Ad

Suggested Title: New Note Edit
29 July 2025 at 12:12 AM

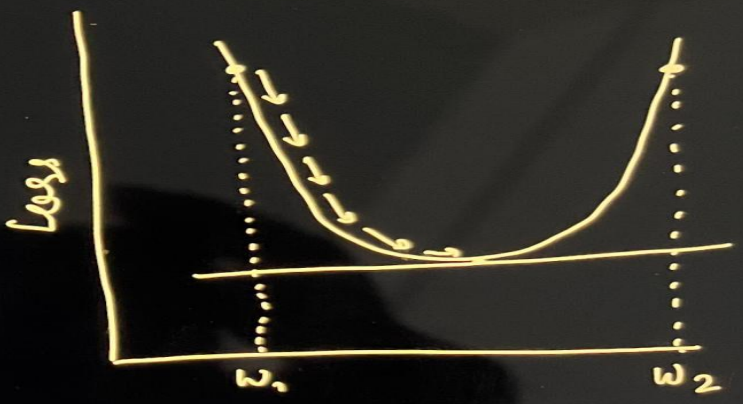
62%

$$w_{\text{new}} = w_{\text{old}} - \eta \frac{\partial L}{\partial w_{\text{old}}}$$

* All the n data points.
 \Rightarrow Gradient Descent.

$$\text{Loss} = \frac{1}{n} \sum_{i=1}^n (y - y')^2$$

* Only 1 point at a time
 \Rightarrow Stochastic Gradient Descent (SGD)

$$\text{Loss} = (y - y')^2$$


Stochastic Gradient Descent

* A sample of 'k' data points.

=> Mini Batch SGD.

$$\text{Loss} = \frac{1}{k} \sum_{i=1}^k (\gamma - \gamma')^2$$



Stochastic Gradient Descent

The image shows handwritten notes on a blackboard. At the top, the title 'Stochastic Gradient Descent' is written. Below it, two gradient vectors are shown: $\left[\frac{\partial L}{\partial W_{old}} \right]$ and $\left[\frac{\partial L}{\partial W_{old}} \right]_{GD}$. The first vector is labeled 'MBSGD' and 'Sample.' with an arrow pointing to it. The second vector is labeled 'GD' and 'population' with an arrow pointing to it. A tilde symbol (\approx) is placed between the two vectors. Below this, a note states: '* The zig-zag movement in SGD is due to the noise.' where 'noise' is circled. The final line of text says: 'which can be dealt with SGD with momentum'.

$$\left[\frac{\partial L}{\partial W_{old}} \right] \approx \left[\frac{\partial L}{\partial W_{old}} \right]_{GD}$$

MBSGD Sample. GD population

* The zig-zag movement in SGD is due to the noise.

which can be dealt with SGD with momentum

Stochastic Gradient Descent

* Exponential weighted moving average :

$$V_t = \beta * V_{t-1} + (1-\beta) \theta_t$$

V_t : Current E.W.M.A. (timestamp)

V_{t-1} : previous E.W.M.A.

θ_t : Current data point

$$V_0 = 0$$

$$V_1 = \cancel{\beta} V_0 + (1-\beta) \theta_1 \\ = (1-\beta) \theta_1$$



Stochastic Gradient Descent

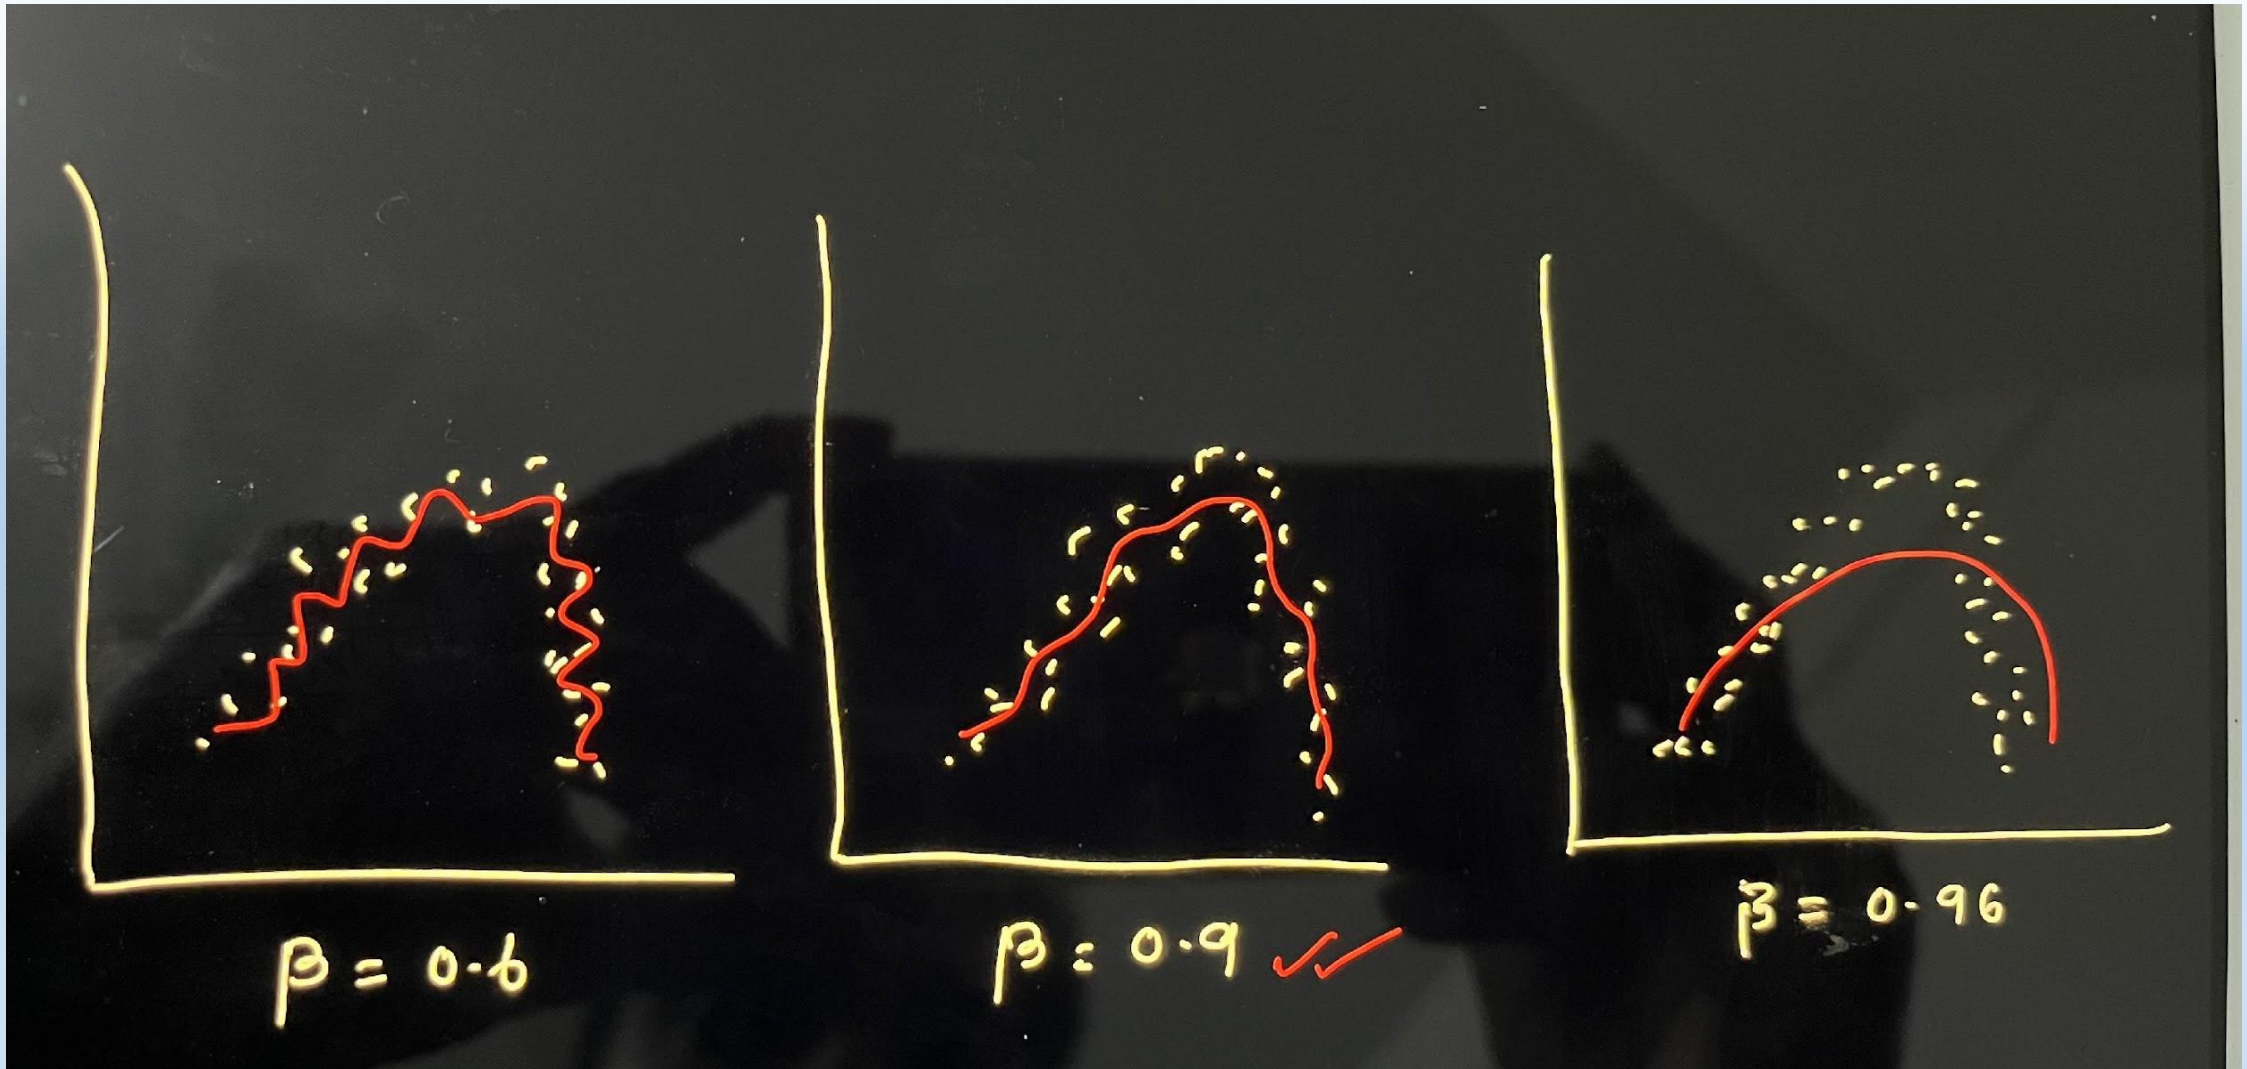
$$\begin{aligned}v_2 &= \beta \cdot v_1 + (1-\beta) \theta_2 \\&= \beta \cdot (1-\beta) \theta_1 + (1-\beta) \theta_2 \\&= (1-\beta) (\beta \cdot \theta_1 + \theta_2)\end{aligned}$$

$$\begin{aligned}v_3 &= \beta \cdot v_2 + (1-\beta) \theta_3 \\&= \beta \cdot (1-\beta) (\beta \cdot \theta_1 + \theta_2) + (1-\beta) \theta_3 \\&= (1-\beta) (\beta^2 \theta_1 + \beta \theta_2 + \theta_3)\end{aligned}$$

$$\vdots$$
$$v_n = (1-\beta) \left(\beta^{n-1} \theta_1 + \beta^{n-2} \theta_2 + \dots + \beta \theta_{n-1} + \theta_n \right)$$

$$0 < \beta < 1$$

Stochastic Gradient Descent



Stochastic Gradient Descent

